

Dineri

Pre-launch readiness review

Next.js 16 App Router · React 19 · Drizzle/PostgreSQL
Audited against the running application, not a checklist.
Generated 2026-09-05

Executive summary

Dineri's SEO foundation is stronger than most pre-launch products. The venue pages - the pages that actually earn search traffic - already ship canonical URLs, Open Graph and Twitter cards, and valid Restaurant structured data with opening hours, social profiles and reservation status. Soft-empty pages are correctly marked noindex, and missing pages return a real 404 rather than a 200.

The gaps are concentrated in two places. First, the marketing site (/, /pricing, /features and the rest) has none of what the venue pages have: no canonical, no Open Graph, no Organization or WebSite structured data. Second, a performance defect affects every page on the platform - 25 font files are preloaded on every single request. Both are fixable in days, not weeks.

One item is a genuine launch blocker: every canonical URL, Open Graph URL and sitemap entry is derived from `BETTER_AUTH_URL`, which falls back to `http://localhost:3000`. If that variable is not set to the public HTTPS domain at deploy time, Google will be told the canonical home of every page is localhost.

Scorecard

Area	Status	Detail
Crawl directives (robots.txt)	PASS	Private areas disallowed; sitemap referenced
XML sitemap	PARTIAL	Dynamic and filtered; no lastModified on marketing, unbounded at scale
Canonical URLs	PARTIAL	Venue pages yes; entire marketing site no
Structured data	PARTIAL	Restaurant schema good; no Organization/WebSite/Menu
Open Graph / Twitter	PARTIAL	Venue pages only; no site-wide default
Indexation control	PASS	noindex on unpublished menus verified
HTTP status codes	PASS	404 returns 404, not soft-200
Titles & descriptions	PASS	Template applied; unique per page
Heading structure	PASS	Exactly one H1 on every page sampled
Core Web Vitals readiness	FAIL	25 font preloads/page; 33 raw tags
Internationalisation	FAIL	Cookie-based locale; no indexable localised URLs
Absolute URL configuration	FAIL	Falls back to localhost if env unset

Method: findings were produced by reading the source and by querying the running application - fetching rendered HTML for /, /pricing, /r/food-mac, /r/food-mac/menu and /preview, inspecting response headers, and requesting known-missing URLs to confirm status codes. Every claim below carries the evidence it came from.

What was measured

Head-tag coverage across representative pages, taken from the rendered HTML of the running app:

Page	Canonical	OG tags	JSON-LD	H1	Status
/	0	0	0	1	200
/pricing	0	0	0	1	200
/preview	0	0	0	1	200
/r/food-mac	1	1	2	1	200
/r/food-mac/menu	1	-	0	1	200
/r/does-not-exist	-	-	-	-	404
/r/code-infortive/menu	-	-	-	-	noindex, follow

The pattern is unambiguous: the venue template is well optimised and the marketing template has had no SEO work at all. The two were clearly built at different times.

Critical - fix before launch

P0-1. Canonical, Open Graph and sitemap URLs can resolve to localhost

P0

Evidence

```
src/lib/seo.ts
const configured =
  process.env.BETTER_AUTH_URL ||
  process.env.NEXT_PUBLIC_BETTER_AUTH_URL ||
  "http://localhost:3000"; <-- silent fallback

.env.example
BETTER_AUTH_URL=http://localhost:3000
```

Why it matters

siteUrl() feeds absoluteUrl(), which feeds every canonical tag, every Open Graph url, and every entry in sitemap.xml. It is derived from an authentication variable, not an SEO one, and it fails silently. If BETTER_AUTH_URL is set to an internal hostname, an http:// origin, or is missing on the deploy target, Google is told the canonical location of every page on the platform is somewhere it cannot reach. This is the single most damaging thing on this list because it invalidates work that is otherwise correct.

Fix

- Introduce a dedicated NEXT_PUBLIC_SITE_URL (for example https://dineri.app) and read it first in siteUrl().
- Throw at build time when NODE_ENV is production and the value is missing or points at localhost - a failed build is far cheaper than a de-indexed launch.
- After deploying, fetch /sitemap.xml and /robots.txt on the live domain and confirm every URL is absolute HTTPS on the public host.
- Decide apex vs www, enforce one with a 301 at the edge, and make sure siteUrl() returns the winner.

Files: src/lib/seo.ts, .env.example, src/app/sitemap.ts, src/app/robots.ts

P0-2. Every page preloads 25 font files

P0

Evidence

```
$ curl -sI http://localhost:3000/ -> 25 rel=preload as=font
$ curl -sI http://localhost:3000/r/food-mac -> 25 rel=preload as=font

src/app/layout.tsx declares 17 next/font families,
all applied to <html> as CSS variables, none with preload:false.
```

Why it matters

The 17 families exist so a merchant can pick a typeface for their venue page. But declaring them in the root layout makes Next preload all of them on every request, including the marketing home page which uses one. Twenty-five render-critical font fetches compete with the LCP image for bandwidth and connections. This directly damages Largest Contentful Paint, which is a confirmed ranking signal, and it is paid on every page view by every visitor on every device.

Fix

- Set preload: false on every family except the one or two the marketing UI actually renders (Inter and Inter Tight).
- Keep the CSS variables - preload:false still ships the @font-face, it just stops the eager fetch, so a venue page that selects Playfair still gets Playfair.
- Better still, load merchant-selectable fonts only on the venue routes rather than in the root layout.
- Re-measure: the same curl should report 2-3 preloads, not 25.

Files: src/app/layout.tsx

P0-3. No canonical tag on any marketing page

P0

Evidence

```
$ curl -s http://localhost:3000/ | grep -c 'rel="canonical"' -> 0
$ curl -s http://localhost:3000/pricing | grep -c 'rel="canonical"' -> 0

Venue pages do have one (src/app/(preview)/r/[slug]/seo.ts):
  alternates: { canonical: path }
```

Why it matters

Without a canonical, every URL variant of a marketing page is a separate candidate for indexing: apex vs www, http vs https, with and without a trailing slash, and - most relevant here - every campaign URL. The platform itself appends `utm_source` and `utm_medium` to QR and share links, so parameterised URLs are guaranteed to exist in the wild. Ranking signals get split across duplicates and Google picks a canonical on your behalf, often the wrong one.

Fix

- Add `alternates: { canonical: '/the-path' }` to the metadata export of every page under (main).
- Extract a small helper so a new marketing page cannot be added without one.
- Set a self-referencing canonical on the home page explicitly - it is the most linked and most parameterised URL you have.

Files: `src/app/(main)/**/*.page.tsx`, `src/lib/seo.ts`

P0-4. The interactive demo is indexable and duplicates the venue template

P0

Evidence

```
$ curl -s http://localhost:3000/preview | grep 'name="robots"' -> (none)

<title>Restaurent - Menu, reservations & links · Dineri</title>
      ^^^^^^^^^^^ typo, and it presents as a real venue
```

Why it matters

`/preview` and its sub-pages render the same layout as a real venue page, for a fictional restaurant, with `noindex`. Google sees a venue page that competes with your customers' genuine venue pages for the same template patterns, on your own domain. The misspelled title would also be the visible snippet.

Fix

- Add `robots: { index: false, follow: true }` to the metadata of `/preview` and every page beneath it.
- Add `/preview` to the disallow list in `robots.ts` as a second line of defence.
- Fix the 'Restaurent' typo regardless - it appears in the browser tab today.

Files: `src/app/(main)/preview/**/*.page.tsx`, `src/app/robots.ts`

High priority

P1-1. No Organization or WebSite structured data

P1

Evidence

```
$ curl -s http://localhost:3000/ | grep -c 'application/ld+json' -> 0
```

Venue pages emit Restaurant JSON-LD correctly; the brand itself emits none.

Why it matters

Organization schema is how Google associates your logo, official name, and social profiles with the brand - the input to a knowledge panel. WebSite schema with SearchAction enables a sitelinks search box. Neither exists, so the brand has no machine-readable identity even though every individual customer venue does.

Fix

- Emit Organization JSON-LD in the root layout: name, url, logo (/images/logo.png), and sameAs with the four social profiles already listed in the footer.
- Emit WebSite JSON-LD with the site name and a potentialAction SearchAction if you add site search.
- Reuse the pattern already proven in r/[slug]/seo.ts rather than inventing a second one.

Files: `src/app/layout.tsx`, `src/lib/seo.ts`

P1-2. No default Open Graph or Twitter card for the marketing site

P1

Evidence

```
$ curl -s http://localhost:3000/ | grep -c 'property="og:' -> 0
```

`src/app/layout.tsx` metadata has title + description only:
`metadataBase`, `title.template`, `description`

Why it matters

Every share of the home page, pricing page or a blog post on WhatsApp, LinkedIn, Slack or X renders as a bare link with no image and whatever text the platform guesses. For a product sold largely through social and word of mouth, this suppresses click-through on exactly the channels you are marketing through. Venue pages already do this properly, so the capability exists.

Fix

- Add openGraph and twitter blocks to the root metadata so every marketing page inherits them.
- Ship a 1200x630 OG image; the existing /images/logo.png is 3172x762 and will be letterboxed badly.
- Override per page where the content justifies it (pricing, features).

Files: `src/app/layout.tsx`, `public/images/`

P1-3. Menu pages carry no structured data

P1

Evidence

```
$ curl -s http://localhost:3000/r/food-mac/menu | grep -c 'application/ld+json' -> 0
```

The venue page links to it: `hasMenu: absoluteUrl(`${path}/menu`)`
...but the menu page itself describes nothing.

Why it matters

This is the largest missed opportunity on the platform. Menu pages hold structured, high-intent data - categories, item names, descriptions, prices, and now dietary and allergen tags - and Google supports Menu, MenuSection and MenuItem schema for exactly this. The allergen work maps directly onto suitableForDiet. Every published menu on the platform is currently invisible as structured content.

Fix

- Emit Menu JSON-LD with a MenuSection per category and a MenuItem per item (name, description, offers.price, offers.priceCurrency).
- Map the dietary tags to schema.org suitableForDiet (VegetarianDiet, VeganDiet, GlutenFreeDiet, HalalDiet) and surface allergens in the item description.
- This scales automatically across every customer venue - build it once in the menu page and every merchant benefits.

Files: `src/app/(preview)/r/[slug]/menu/page.tsx`, `src/utlis/tags.ts`

P1-4. 33 raw `<img>` tags against 2 uses of `next/image`

P1

Evidence

```
$ grep -rn '<img' src/ --include=*.tsx | wc -l    -> 33
$ grep -rl 'from "next/image"' src/ | wc -l     -> 2
```

Includes venue background photos, menu item photos and gallery images.

Why it matters

Raw `<img>` without width and height causes layout shift as images arrive, and Cumulative Layout Shift is a Core Web Vital. It also means no responsive srcset, no AVIF/WebP conversion, and full-resolution originals served to phones. Menu and venue pages are image-dense and mobile-first - this is where the traffic is.

Fix

- Set explicit width and height (or an aspect-ratio container) on every `img` to reserve space, even where `next/image` is impractical.
- Move merchant uploads to `next/image` with `remotePatterns` for the S3 bucket, which the CSP already allows.
- Add `fetchPriority="high"` to the venue page LCP image and keep `loading="lazy"` on the rest.
- Note: the venue background layer was deliberately left as a plain `img` (arbitrary merchant hosts); give it width/height rather than converting it.

Files: `src/app/(preview)/r/[slug]**`, `src/components/shared/appearance-background.tsx`

P1-5. Six locales exist but none can be indexed

P1

Evidence

```
src/i18n/request.ts
  const locale = store.get("locale")?.value || "en";

messages/: de.json en.json es.json fr.json it.json nl.json
No /[locale] route segment, no hreflang tags emitted.
```

Why it matters

Locale is chosen by cookie, so every language shares one URL. Googlebot crawls without cookies, sees only English, and has no alternate URL to index. Five translations are therefore invisible to search. If the product targets Dutch, German, French, Italian or Spanish restaurants, this caps organic reach in those markets at zero regardless of content quality.

Fix

- If localised SEO matters: move to `/[locale]` path segments with `next-intl` routing and emit `hreflang` plus `x-default` via `alternates.languages`.

- If it does not matter yet: leave as is, but keep the cookie approach out of the indexable surface and revisit before entering a non-English market.
- Either way this is a routing decision, not a copy decision - worth settling before launch rather than after URLs are indexed.

Files: `src/i18n/request.ts`, `next.config.ts`

Medium priority

#	Issue	Sev	Action
M1	Sitemap marketing entries have no lastModified	P2	Add dates so Google can prioritise recrawls
M2	sitemap.ts selects every restaurant with no limit	P2	Paginate into a sitemap index before 50k URLs / 50MB
M3	No breadcrumb structured data on venue sub-pages	P2	Add BreadcrumbList to /menu, /reserve, /track-order
M4	No icons or web manifest in metadata	P2	Add icons + apple-touch-icon + manifest.webmanifest
M5	/preview not listed in robots.ts disallow	P2	Add alongside the noindex from P0-4
M6	Marketing pages lack per-page OG overrides	P2	Follow after P1-2 for pricing and features
M7	Venue pages are client components	P2	Content is server-rendered (verified) but hydration is heavy; consider RSC for static parts
M8	Production caching unverified	P2	Dev sends no-cache; confirm ISR/CDN headers on the real deploy
M9	No XML sitemap for images	P2	Optional; useful given how image-led venue pages are
M10	Root description claims '12,400+ restaurants'	P2	Unverified claim in the meta description - confirm or soften

What is already right

These are working and should not be disturbed. They were verified against the running application, not assumed.

- robots.txt disallows /dashboard, /admin, /onboarding, /api/, the auth routes and the order/reservation confirmation pages, and points at the sitemap.
- sitemap.xml is generated from the database, revalidates hourly, and deliberately omits venue menu and reserve pages that a visitor could not use - matching the noindex applied to those same pages.
- Venue pages emit a canonical, Open Graph and Twitter tags, and Restaurant JSON-LD including openingHoursSpecification, sameAs, hasMenu and acceptsReservations.
- An unpublished menu returns 'noindex, follow' - confirmed on /r/code-infotve/menu.
- Missing pages return HTTP 404, not a soft 200 - confirmed on two different URL shapes.
- Title template '%s · Dineri' is applied consistently and descriptions are unique per page.
- Exactly one H1 on every page sampled.
- Security headers including a real Content-Security-Policy are configured in next.config.ts.

Recommended sequence

Ordered by damage prevented per hour of work. The first block is genuinely blocking; the rest can land in the first weeks after launch without harm.

When	Do	Why then
Before launch	P0-1 site URL env + build-time guard	Wrong canonicals are hard to undo once crawled
Before launch	P0-3 canonicals on marketing pages	Must exist before the first crawl, not after
Before launch	P0-4 noindex /preview	Stops a duplicate entering the index at all
Before launch	P0-2 font preloads	One-line change, large CWV win, no risk
Launch week	P1-1 Organization + WebSite JSON-LD	Brand identity starts accruing from day one
Launch week	P1-2 default OG image and tags	Every launch share benefits immediately
First month	P1-3 Menu structured data	Highest ceiling; needs care to get right
First month	P1-4 image dimensions and next/image	Steady CWV improvement across all venues
Decide early	P1-5 localisation routing	Cheap now, expensive once URLs are indexed
Ongoing	M1-M10	Housekeeping

Verify after deploying

Run these against the live domain, not localhost. Each maps to a finding above.

```
curl -sI https://YOUR-DOMAIN/ | grep -c 'as="font"' # expect 2-3, not 25
curl -s https://YOUR-DOMAIN/ | grep -c 'rel="canonical"' # expect 1
curl -s https://YOUR-DOMAIN/ | grep -c 'og:' # expect > 0
curl -s https://YOUR-DOMAIN/preview | grep 'name="robots"' # expect noindex
curl -s https://YOUR-DOMAIN/sitemap.xml | grep -c 'localhost' # expect 0
curl -s https://YOUR-DOMAIN/robots.txt | grep Sitemap # expect https URL
```

Then: submit the sitemap in Google Search Console, request indexing for the home page, and validate one venue page plus one menu page in the Rich Results Test.

Scope and limits of this audit

This audit was run against the development server on localhost. Findings about markup, structured data, status codes and configuration transfer directly to production because they come from the same code. Findings about performance are directional rather than absolute: a development build is unoptimised, so treat the 25 font preloads as a structural fact (it is caused by the layout, not the build) but do not read the page weights as production numbers.

Not covered: off-page factors (backlinks, brand mentions), content strategy and keyword targeting, log-file analysis, real Core Web Vitals field data from CrUX, and anything requiring Search Console access. Lighthouse was not run - the numbers here come from response headers and rendered HTML, which are reproducible.